FRANCISCO PERDIGON ROMERO

THE REAL TALK

6 things every AI engineer should know before shipping RAG to production

Swipe to see tip 1 →

FRANCISCO PERDIGON ROMERO

TIP 1

Retrieval quality beats model size every time

- → Chunk strategy defines what the retriever finds
- → Bad chunks = hallucinated answers, every time
- → I test chunking before I touch the LLM

Swipe to see tip 2 →

FRANCISCO PERDIGON ROMERO

TIP 2

Agent boundaries prevent cascade failures in production

- → Each agent owns one job — no cross-agent side effects
- → I define explicit input/output contracts per agent
- → Failures stay local; they don't cascade

Swipe to see tip 3 →

FRANCISCO PERDIGON ROMERO

TIP 3

Observability is not optional past the prototype stage

- → Log every retrieval call with query, chunks, scores
- → Trace agent decisions end-to-end before going live
- → You can't debug what you can't see

Swipe to see tip 4 →

FRANCISCO PERDIGON ROMERO

TIP 4

Evaluate retrieval and generation separately

- → Retrieval eval: does the right chunk appear in top-k?
- → Generation eval: is the answer grounded in context?
- → Combined metrics hide which layer is failing

Swipe to see tip 5 →

FRANCISCO PERDIGON ROMERO

TIP 5

Guard against prompt injection at every agent boundary

- Retrieved docs can carry adversarial instructions
- I sanitize and scope context before passing to LLM
- One unguarded boundary can hijack the whole pipeline

Swipe to see tip 6 →

FRANCISCO PERDIGON ROMERO

THE MOST IMPORTANT

Start with one agent. Earn complexity.

- → Most failed RAG deployments were over-architected early
- → I shipped a single-agent pipeline first; scaled later
- → Complexity is a liability until the baseline works

Swipe for the CTA →

FOLLOW

Follow Francisco Perdigon Romero

I post on agentic AI, RAG pipelines, and ML for telecom — applied lessons from production systems.

linkedin.com/in/fperdigon